

Hidden keys and repeated biometric records

Can multiple protected records reveal identity when keys remain hidden?

Fresh-key learned attacks: uncertainty remains; no equivalence claim.


Recurring transforms: large gains from multiple same-identity records.

Boundary location changes with the dataset and identity partition.

Four datasets; hidden-key protocols; matched MOBIO/FEI follow-up with 216 trained endpoints.


Experimental architecture

A DATA AND PROTECTION



B KEY-BLIND ATTACK AND EVALUATION

Attacker observes protected records, not keys or slot labels; slots are revealed only in a labelled control.



KEY DESIGN Fresh: source-record keys are split-disjoint. Pool K: hidden transforms recur across identity splits.

REPLICATION MOBIO / FEI follow-up: 2 identity partitions x 3 model seeds; matched 120-epoch training caps.

DESIGN Same-identity aggregation under controlled transform reuse; detailed attacker paths in the companion figure.

Train on separate identities. Reconstruct an embedding, then link to a held-out gallery.

Dataset and protection coverage

Dataset	Identities	Train/val/test	Embeddings	Variation
MOBIO	150	90/30/30	1,799 / 1,800	Sessions
LFW	125	75/25/25	1,500 / 1,500	In-the-wild
FEI	200	120/40/40	2,378 / 2,400	Pose/expression
SCface	130	78/26/26	2,851 / 2,860	Camera/distance

BioHash: 128 bits; four datasets; includes a Haar-sign control on MOBIO.
MLP-Hash: 512 bits; MOBIO; paper-specified, not source-exact.
IoM-GRP: 300 codes (q=16); PolyProtect: 170 reals (m=5, overlap=2).
Both: MOBIO/FEI three-seed, two-partition follow-up; SCface one-seed pilots.

SCface: mugshot gallery and visible surveillance probes; 9 detection failures, all identities eligible.

Set reconstruction and gallery linkage


A MATCHED EXPOSURE CONSTRUCTION

One identity, n source records; nested subsets from one fixed permutation per repeat; gallery excluded.

Fresh: independent source-record keys. Pool K: hidden transforms recur across identity splits. No keys or slots in model input.

T: $B \times n \times d$ $d = 128$ BioHash; 512 MLP-Hash; 4,800 one-hot IoM-GRP; 170 PolyProtect

B TWO AGGREGATION PATHS



Primary: mean-pool MLP. Secondary: DeepSets. m masks valid records; all n records are valid in these runs.

C SUPERVISION

Target $u = \text{normalize}(\text{mean of exposed source embeddings})$.

Loss = $\text{mean}[1 - \cos(z, u)] + 0.1 \times \text{MSE}(z, u)$.

Adam: lr 0.001; weight decay 0.0001; hidden width 256.

Train identities only; select minimum validation loss; freeze test.

D HELD-OUT IDENTITY LINKAGE

Scores = $z \cdot G^T$; G: $N \times 512$ unprotected unit gallery.

Rank identities; top-1/top-5, AUROC, EER, TAR at FAR.


Follow-up: 3 model seeds x 2 identity partitions; 120 epochs.

Crossed seed/identity intervals; paired primary contrasts.

B: batch size; n: exposure count; d: template input width; N: test gallery identities. Source targets are not inference inputs.

Two aggregation paths; shared training objective; held-out identity linkage. Key values remain hidden.


Earlier three-seed key-pool studies



73 conditions from 12 earlier studies. New one-seed pilots are reported separately.


Mechanism and correlation controls

MOBIO mechanism controls (30 gallery identities; chance 3.33%)



Slot identifiers are not key values. Coarse and fine correlation sweeps use separate partitions.



Replicated multi-record amplification



Bars: crossed model-seed / identity bootstrap 95% intervals. Tests: identity sign-flips, Holm family of 8.
Matched 120-epoch caps; key/set seeds fixed. Partitions are separate sensitivity analyses, not independent replications.

MOBIO / FEI; IoM-GRP / PolyProtect; 3 model seeds x 2 identity partitions; matched 120-epoch caps.

Native linkage and radial sensitivity



Native tests: 12 / 12 Holm-adjusted $p \leq 0.006$; identity-clustered 95% intervals; dashed lines: chance.

A/B: splits 91831/91843. Key seeds 1/2/3: 91873/91879/91883. Gallery-label permutations preserve probe clusters.

Radial stress: 32 held-out records per dataset/partition, fixed keys; measures scale sensitivity, not natural norm leakage.

Protected-gallery matching is distinct from learned linkage. Radial sensitivity does not establish natural norm leakage.